

Lecture 6

Regression for completely randomized experiment



Outline

- Using regression with no covariates
- Using regression with covariates adjustments
- Using regression with covariates adjustments and interactions
- The LRC-CPPT cholesterol data example

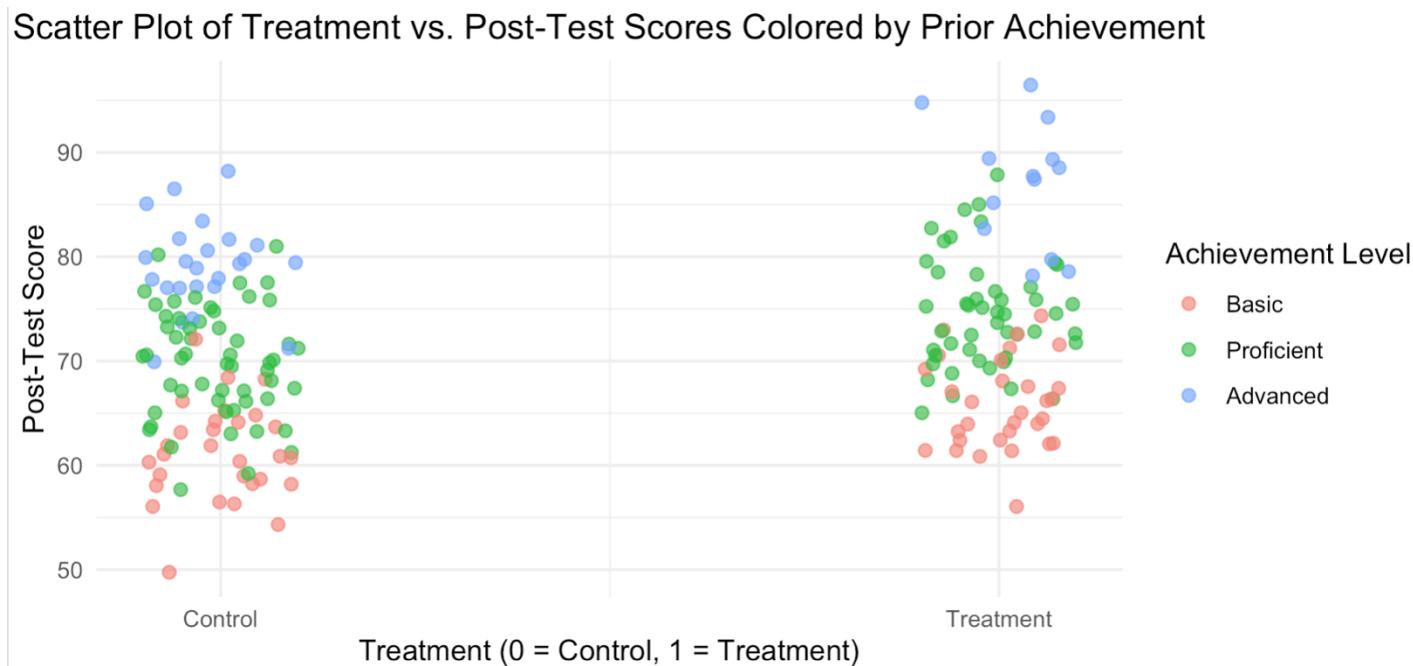
Linear regression and causality

- Linear regression:

$$\mathbb{E}(Y_i|W_i, X_i) = \alpha + \gamma W_i + \beta^T X_i$$

- Benefits of using linear regression:

- Adjust for confounding variables
 - Not need for completely randomized experiments as pre-treatment covariates are not confounded
- More accurate estimator if covariates explain part of the noise in the outcome



Linear regression and causality

- Linear regression:

$$\mathbb{E}(Y_i|W_i, \mathbf{X}_i) = \alpha + \gamma W_i + \boldsymbol{\beta}^T \mathbf{X}_i$$

- Question:

- When can we interpret the coefficient(s) as causal effect?
- How can we do correct inference if we take into account the randomization procedure of treatment assignments?

- Some critiques

- In completely randomized experiments, covariates are not confounders
- Why do we want to assume a linear model if we don't need to?
 - Model $\mathbb{E}(Y_i|W_i, \mathbf{X}_i) = \alpha + \gamma W_i + \boldsymbol{\beta}^T \mathbf{X}_i$ assumes same causal effect for all levels of \mathbf{X}_i

“Experiments should be analyzed as experiments, not as observational studies”

---- David A. Freedman, 2006

The LRC-CPPT cholesterol data

- An experiment to evaluate the effect of the drug cholestyramine on reducing cholesterol levels
- $N = 337$ patients are completely randomized
- **Pre-treatment covariates:** two cholesterol measurements before and after a suggestion of low-cholesterol diet, both measurements taken prior to the random assignment
 - $\text{cholp} = 0.25 \text{ chol1} + 0.75 \text{ chol2}$

Table 7.1. *Summary Statistics for PRC-CPPT Cholesterol Data*

	Variable	Control ($N_c = 172$)		Treatment ($N_t = 165$)		Min	Max
		Average	Sample (S.D.)	Average	Sample (S.D.)		
Pre-treatment	chol1	297.1	(23.1)	297.0	(20.4)	247.0	442.0
	chol2	289.2	(24.1)	287.4	(21.4)	224.0	435.0
	cholp	291.2	(23.2)	289.9	(20.4)	233.0	436.8
Post-treatment	cholf	282.7	(24.9)	256.5	(26.2)	167.0	427.0
	chold	-8.5	(10.8)	-33.4	(21.3)	-113.3	29.5
	comp	74.5	(21.0)	59.9	(24.4)	0	101.0

The LRC-CPPT cholesterol data

- An experiment to evaluate the effect of the drug cholestyramine on reducing cholesterol levels
- $N = 337$ patients are completely randomized
- **Post-treatment outcomes:**
 - cholf: post-treatment average cholesterol level
 - chold = cholf – cholp
 - comp: compliance rate, the percentage of individuals follow the treatment assignment

Table 7.1. *Summary Statistics for PRC-CPPT Cholesterol Data*

	Variable	Control ($N_c = 172$)		Treatment ($N_t = 165$)		Min	Max
		Average	Sample (S.D.)	Average	Sample (S.D.)		
Pre-treatment	chol1	297.1	(23.1)	297.0	(20.4)	247.0	442.0
	chol2	289.2	(24.1)	287.4	(21.4)	224.0	435.0
	cholp	291.2	(23.2)	289.9	(20.4)	233.0	436.8
Post-treatment	cholf	282.7	(24.9)	256.5	(26.2)	167.0	427.0
	chold	-8.5	(10.8)	-33.4	(21.3)	-113.3	29.5
	comp	74.5	(21.0)	59.9	(24.4)	0	101.0

The LRC-CPPT cholesterol data

- Can we evaluate the drug effect by simply look at whether chold is positive or negative?
 - No! The before-after comparison is NOT necessarily causal
 - Even for the control group, chold is significantly negative
- The patient's post-treatment cholesterol should be highly correlated with his/her pre-treatment cholesterol level
- How do we evaluate the causal effect after “adjusting for the pre-treatment cholesterol”?
 - Adjust for pre-treatment cholesterol by regression

Table 7.1. Summary Statistics for PRC-CPPT Cholesterol Data

	Variable	Control ($N_c = 172$)		Treatment ($N_t = 165$)			
		Average	Sample (S.D.)	Average	Sample (S.D.)	Min	Max
Pre-treatment	chol1	297.1	(23.1)	297.0	(20.4)	247.0	442.0
	chol2	289.2	(24.1)	287.4	(21.4)	224.0	435.0
	cholp	291.2	(23.2)	289.9	(20.4)	233.0	436.8
Post-treatment	cholf	282.7	(24.9)	256.5	(26.2)	167.0	427.0
	chold	-8.5	(10.8)	-33.4	(21.3)	-113.3	29.5
	comp	74.5	(21.0)	59.9	(24.4)	0	101.0

Linear regression with no covariates

- Neyman's approach

$$\hat{\tau}^{\text{dif}} = \bar{Y}_t^{\text{obs}} - \bar{Y}_c^{\text{obs}}$$

$$\bar{Y}_c^{\text{obs}} = \frac{1}{N_c} \sum_{i:W_i=0} Y_i^{\text{obs}} \quad \text{and} \quad \bar{Y}_t^{\text{obs}} = \frac{1}{N_t} \sum_{i:W_i=1} Y_i^{\text{obs}}$$

- An alternative way to get $\hat{\tau}^{\text{dif}}$ is by linear regression

$$(\hat{a}, \hat{b}) = \arg \min_{(a,b)} \sum_{i=1}^N (Y_i - a - bW_i)^2$$

- It is easy to show that $\hat{b} = \hat{\tau}^{\text{dif}}$

- $\arg \min_{(a,b)} \sum_{i=1}^n (Y_i - a - bW_i)^2 = \operatorname{argmin}_{(a,b)} [\sum_{i:W_i=0} (Y_i - a)^2 + \sum_{i:W_i=1} (Y_i - a - b)^2]$
- $\hat{a} = \bar{Y}_c^{\text{obs}}, \hat{a} + \hat{b} = \bar{Y}_t^{\text{obs}}$
- However, if we use R function `lm()`, the variance estimator is $\frac{1}{N} \frac{(N_c-1)s_c^2 + (N_t-1)s_t^2}{N-2}$
 - Different from $\frac{s_c^2}{N_c} + \frac{s_t^2}{N_t}$ in Neyman's approach

Causal interpretation of the linear model

- Linear model on the potential outcomes

$$Y_i(w) = \alpha + \tau_i w + \varepsilon_i^* = \alpha + \tau w + \varepsilon_i(w)$$

where $\mathbb{E}(\varepsilon_i^*) = 0$ and $\varepsilon_i(w) = \varepsilon_i^* + (\tau_i - \tau)w$

- Not really an assumption if w only has two values

$$Y_i(0) = \alpha + \varepsilon_i^*, Y_i(1) = Y_i(0) + \tau_i$$

- Assume that there is a super-population and the potential outcomes are i.i.d. samples
- The observed outcomes $Y_i = W_i Y_i(1) + (1 - W_i) Y_i(0)$ are **not i.i.d. samples** under complete randomization
- Define PATE: $\tau = \mathbb{E}(\tau_i) = \mathbb{E}(Y_i(1) - Y_i(0))$
- $\alpha = \mathbb{E}(Y_i(0))$ and $\mathbb{E}(\varepsilon_i(w)) = 0$

Linear regression with no covariates

- Causal model on the potential outcomes

$$Y_i(w) = \alpha + \tau_i w + \varepsilon_i^* = \alpha + \tau w + \varepsilon_i(w)$$

where $\mathbb{E}(\varepsilon_i^*) = 0$ and $\varepsilon_i(w) = \varepsilon_i^* + (\tau_i - \tau)w$

- If the treatment is binary ($w = 0, 1$), then the above model essentially has no assumption on $Y_i(0)$ and $Y_i(1)$
- If the treatment is continuous, the model assumes a linear but heterogeneous causal effect on each individual
- How to estimate τ from observed data?
- When does the above model imply the linear regression model on observed data?

$$Y_i^{\text{obs}} = \alpha + \tau W_i + \varepsilon_i$$

Linear regression with no covariates

$$Y_i(w) = \alpha + \tau_i w + \varepsilon_i^* = \alpha + \tau w + \varepsilon_i(w)$$

We assume the following identification conditions

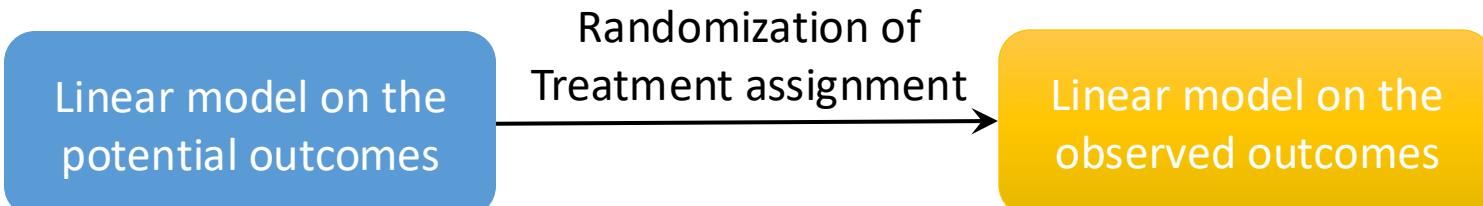
- **Randomization of the treatment:**

$$(\mathbf{Y}(0), \mathbf{Y}(1)) \perp \mathbf{W}$$

- Satisfied in completely randomized experiments
- Then, $\mathbb{E}(Y_i(w)) = \mathbb{E}(Y_i(w)|W_i = w) = \mathbb{E}(Y_i^{\text{obs}}|W_i = w) = \alpha + \tau w$
- So this implies a regression model

$$Y_i^{\text{obs}} = \alpha + \tau W_i + \varepsilon_i$$

$$\text{where } \varepsilon_i = \varepsilon_i(W_i) = \varepsilon_i^* + (\tau_i - \tau)W_i$$



- What is the correct statistical inference?

Linear regression with no covariates

- regression model

$$Y_i^{\text{obs}} = \alpha + \tau W_i + \varepsilon_i$$

where $\varepsilon_i = \varepsilon_i(W_i) = \varepsilon_i^* + (\tau_i - \tau)W_i$

- Follow the linear regression convention, we perform statistical inference conditional on (W_1, \dots, W_N)
 - we treat assignment vectors as fixed
- Random sampling of the units
 - $(\varepsilon_i(0), \varepsilon_i(1))$ are independent across i
 - This implies that ε_i in the linear regression model are **independent** as W_i are treated as fixed
 - But they may not follow the same distribution

Homoscedastic error assumption

Homoscedastic error assumption: $\mathbb{V}(\varepsilon_i(0)) = \mathbb{V}(\varepsilon_i(1)) = \sigma^2$

- Then $\mathbb{V}(Y_i^{\text{obs}} | W_i) = \varepsilon_i = \varepsilon_i(W_i)$ always has variance σ^2
- Under homoscedasticity, OLS estimates of the variance is

$$\hat{\sigma}_{Y|W}^2 = \frac{1}{N-2} \sum_{i=1}^N \hat{\varepsilon}_i^2 = \frac{1}{N-2} \sum_{i=1}^N (Y_i^{\text{obs}} - \hat{Y}_i^{\text{obs}})^2,$$

where the estimated residual is $\hat{\varepsilon}_i = Y_i^{\text{obs}} - \hat{Y}_i^{\text{obs}}$, and the predicted value \hat{Y}_i^{obs} is

$$\hat{Y}_i^{\text{obs}} = \begin{cases} \hat{\alpha}^{\text{ols}} & \text{if } W_i = 0, \\ \hat{\alpha}^{\text{ols}} + \hat{\tau}^{\text{ols}} & \text{if } W_i = 1. \end{cases}$$

- Same as the standard linear regression approach

Heteroscedastic errors

- If we don't want to assume $\mathbb{V}(\varepsilon_i(0)) = \mathbb{V}(\varepsilon_i(1))$, then the homoscedastic error assumption fails
 - ε_i has the same distribution for $W_i = 0$, and the same distribution for $W_i = 1$
 - We should use same variance within the treated and control group
 - That leads to the variance estimator in Neyman's approach
- This is also called the Sandwich estimator that is robust to the violation of the homoscedastic noise assumption in linear regression
 - In R, it corresponds to Sandwich estimator with HC2 adjustment

Linear regression with no covariates

To summarize the logic

- We build a (linear) model on the potential outcomes
- This model implies a linear regression model on the observed outcome if $(Y(0), Y(1)) \perp W$
- The coefficient on W_i in the linear regression model is the average causal effect (PATE)
- The linear regression model treat W as fixed so it works for any randomization assignment mechanism that satisfy $(Y(0), Y(1)) \perp W$
- Noise in the linear regression model are independent as long as potential outcomes are independent across units
- For statistical inference
 - The OLS estimator estimator is always unbiased
 - We can apply standard linear regression inference results if we assume $V(\varepsilon_i(0)) = V(\varepsilon_i(1))$
 - If $V(\varepsilon_i(0)) \neq V(\varepsilon_i(1))$, we need to use the robust variance estimator

Linear regression with covariates adjustment

- What are model assumption on the potential outcomes that lead to a linear model on the observed outcome
- Assumption 1: $\mathbb{E}(Y_i(0) | \mathbf{X}_i) = \alpha + \boldsymbol{\beta}^T \mathbf{X}_i$
- Assumption 2: CATE $\tau(\mathbf{x}) = \mathbb{E}(\tau_i | \mathbf{X}_i = \mathbf{x}) \equiv \tau = \text{PATE}$ constant across levels of \mathbf{X}_i
 - We can allow for heterogeneous causal effect but need $\mathbb{E}(\tau_i - \tau | \mathbf{X}_i) = 0$
(individual causal effects are independent from the pre-treatment covariates)
- Then $\mathbb{E}(Y_i(w) | \mathbf{X}_i) = \mathbb{E}(Y_i(0) + \tau_i w | \mathbf{X}_i) = \alpha + \tau w + \boldsymbol{\beta}^T \mathbf{X}_i$
- Unconfoundedness property:
$$(\mathbf{Y}(0), \mathbf{Y}(1)) \perp \mathbf{W} | \mathbf{X}$$
 - $\mathbb{E}(Y_i^{\text{obs}} | W_i = w, \mathbf{X}_i = \mathbf{x}) = \mathbb{E}(Y_i(w) | \mathbf{X}_i = \mathbf{x}) = \alpha + \tau w + \boldsymbol{\beta}^T \mathbf{X}_i$
 - Statistical inference is conditional on both \mathbf{X}_i and W_i

OLS with covariates adjustment

$$(\hat{\alpha}^{\text{ols}}, \hat{\tau}^{\text{ols}}, \hat{\beta}^{\text{ols}}) = \arg \min_{\alpha, \tau, \beta} \sum_{i=1}^N \left(Y_i^{\text{obs}} - \alpha - \tau \cdot W_i - X_i \beta \right)^2$$

- The estimator $\hat{\tau}^{\text{ols}}$ is unbiased for the causal estimand τ
- Even if the model is **incorrect** (either the violation of $\mathbb{E}(Y_i(0)|X_i) = \alpha + \boldsymbol{\beta}^T \mathbf{X}_i$ or $\tau \equiv \mathbb{E}(\tau_i|X_i = \mathbf{x})$), $\hat{\tau}^{\text{ols}}$ still converges to the PATE $\mathbb{E}(\tau_i)$ under complete randomization

Efficiency gain from regression

- If the model is correct, we have

$$\mathbb{V}(\hat{\tau}^{\text{ols}}) \approx \frac{\mathbb{E}\{\mathbb{V}(Y_i(1)|X_i)\}}{N_t} + \frac{\mathbb{E}\{\mathbb{V}(Y_i(0)|X_i)\}}{N_c} \leq \frac{\sigma_c^2}{N_c} + \frac{\sigma_t^2}{N_t}$$

- If X_i is predictive of the (potential) outcomes, we have a more accurate estimator
- If the linear model is **incorrect**, the efficiency might be lost
(Freedman 2008, *Adv. Appl. Math.*)

Estimate of the variance of $\hat{\tau}^{\text{ols}}$ with covariates adjustment

- Assume homoscedastic error assumption:

$$\mathbb{V}(\varepsilon_i(0)) = \mathbb{V}(\varepsilon_i(1)) = \sigma^2 = \mathbb{V}(Y_i^{\text{obs}} | W_i, X_i)$$

We can follow standard linear regression inference and estimate variance of $\hat{\tau}^{\text{ols}}$ as

$$\hat{\mathbb{V}}_{\text{sp}}^{\text{homo}} = \frac{1}{N(N - 1 - \dim(X_i))} \cdot \frac{\sum_{i=1}^N (Y_i^{\text{obs}} - \hat{\alpha}^{\text{ols}} - \hat{\tau}^{\text{ols}} - X_i \hat{\beta}^{\text{ols}})^2}{\bar{W} \cdot (1 - \bar{W})}$$

- The robust variance estimator (Sandwich estimator) without assuming homoscedasticity

$$\begin{aligned}\hat{\mathbb{V}}_{\text{sp}}^{\text{hetero}} &= \frac{1}{N(N - 1 - \dim(X_i))} \\ &\cdot \frac{\sum_{i=1}^N (W_i - \bar{W})^2 \cdot (Y_i^{\text{obs}} - \hat{\alpha}^{\text{ols}} - \hat{\tau}^{\text{ols}} - X_i \hat{\beta}^{\text{ols}})^2}{(\bar{W} \cdot (1 - \bar{W}))^2}\end{aligned}$$

Linear regression with covariates adjustment and interactions

What if the assumption $\tau \equiv \tau(\mathbf{x}) = \mathbb{E}(\tau_i | \mathbf{X}_i = \mathbf{x})$ constant across levels of \mathbf{X}_i is incorrect?

- Assume CATE $\tau(\mathbf{x}) = \mathbb{E}(\tau_i | \mathbf{X}_i = \mathbf{x}) = \tau + \boldsymbol{\gamma}^T(\mathbf{x} - \bar{\mathbf{X}})$
 - τ is still the population average treatment effect
- Still assume $\mathbb{E}(Y_i(0) | \mathbf{X}_i) = \alpha + \boldsymbol{\beta}^T \mathbf{X}_i$
- Then $\mathbb{E}(Y_i(w) | \mathbf{X}_i) = \mathbb{E}(Y_i(0) + \tau_i w | \mathbf{X}_i) = \alpha + \tau w + \boldsymbol{\beta}^T \mathbf{X}_i + \boldsymbol{\gamma}^T(\mathbf{X}_i - \bar{\mathbf{X}})w$
- When does the above model imply the linear regression model with interactions on observed data?

$$Y_i^{\text{obs}} = \alpha + \tau W_i + \boldsymbol{\beta}^T \mathbf{X}_i + \boldsymbol{\gamma}^T(\mathbf{X}_i - \bar{\mathbf{X}})W_i + \varepsilon_i$$

- Unconfoundedness property \rightarrow check by yourself
- In completely randomized experiments, with the interaction terms, we can always guarantee no efficiency loss even when the linear model is wrong (Peng's book section 6.2.2)

Results on the LRC-CPPT cholesterol data

- We estimate the PATE for both the post-treatment cholesterol level cholf and compliance
 - A considerable reduction of the variance of $\hat{\tau}^{\text{ols}}$ for cholf when we add the pre-treatment cholesterol levels in the regression
 - Our goal is always estimating PATE even after “covariates adjustment”
 - In randomized experiments satisfying $(Y(0), Y(1)) \perp W$, adjusting for covariates or not, our estimate of PATE is always valid, we only change the efficiency of our estimate

Covariates	Effect of Assignment to Treatment on			
	Post-Cholesterol Level		Compliance	
	$\hat{\tau}$	(s. e.)	$\hat{\tau}$	(s. e.)
No covariates	-26.22	(3.93)	-14.64	(3.51)
cholp	-25.01	(2.60)	-14.68	(3.51)
chol1, chol2	-25.02	(2.59)	-14.95	(3.50)
chol1, chol2, interacted with W	-25.04	(2.56)	-14.94	(3.49)

The LRC-CPPT cholesterol data

A bit explanation about compliance

- If we compare between control and treatment group, we are evaluating the causal effect of “being assigned”, not the causal effect of actually taking the drug
- Compliance lower in the treatment group possibly due to the side effect of the drug
- Can we just throw away individuals who do not follow the treatment and estimate the causal effect of taking the drug based on the rest individuals? **No**
- Will discuss more about compliance in later lectures

Table 7.1. *Summary Statistics for PRC-CPPT Cholesterol Data*

	Variable	Control ($N_c = 172$)		Treatment ($N_t = 165$)			
		Average	Sample (S.D.)	Average	Sample (S.D.)	Min	Max
Pre-treatment	chol1	297.1	(23.1)	297.0	(20.4)	247.0	442.0
	chol2	289.2	(24.1)	287.4	(21.4)	224.0	435.0
	cholp	291.2	(23.2)	289.9	(20.4)	233.0	436.8
Post-treatment	cholf	282.7	(24.9)	256.5	(26.2)	167.0	427.0
	chold	-8.5	(10.8)	-33.4	(21.3)	-113.3	29.5
	comp	74.5	(21.0)	59.9	(24.4)	0	101.0

Why do we use linear regression in randomized experiments?

- Covariate adjustment can be used to improve efficiency in randomized experiments
 - Always add interaction terms (between each covariate and treatment) to guarantee power improvement
- In completely randomized experiments
 - No need to worry about model misspecification
 - Treatment and covariates are independent